

# A New Approach for Parallel Functional Arrays

Ananya Kumar

Carnegie Mellon University  
ananyak@andrew.cmu.edu

Guy Blleloch

Carnegie Mellon University  
blleloch@cs.cmu.edu

## Abstract

In this paper we introduce a  $O(1)$  wait-free, parallel, functional array, that allows  $O(1)$  reads and writes to the most recent version of the array. We describe the cost dynamics and sketch out a provable implementation. We show favorable benchmarks comparing our functional arrays with regular arrays in Java.

**Categories and Subject Descriptors** CR-number [subcategory]: third-level

**Keywords** array, parallel, cost semantics

## 1. Introduction

Arrays are very important in functional programming languages because they allow work-efficient implementations of algorithms like depth-first search. Accessing old versions of arrays can be useful for efficient checkpointing, logging, and event handling.

### 1.1 Previous Approaches

There has been a lot of great work on functional arrays, but previous approaches have some limitations. Many functional programming languages use monads to build arrays. However, monads do not allow accessing old versions of an array.

Another creative approach is to use compiler based reference counting. If the number of references to an array is provably one, the compiler allows the array to mutate, otherwise the compiler copies the array before writing to the array. However, this approach makes it difficult for programmers to reason about the time complexity of array operations because it depends on the compiler. Furthermore, if multiple variables reference the same array, the array is copied even if all writes are to the most recent version. This could make the time complexity of programs significantly higher than they need to be.

A third approach is to use linear types to ensure that programmers can only access and mutate arrays in valid ways. This approach is very difficult to implement in existing programming languages.

O' Neill's functional arrays guarantee  $O(1)$  access and insertions when used like an imperative array. However, the arrays explicitly implement version trees which are inefficient in practice. Additionally, the reads and writes to the most recent version of the array might take  $O(\log n)$  work if the array is not used like an imperative array.

Chuang's approach supports  $O(1)$  accesses and insertions to the most recent version of arrays, however accessing old versions of the array takes  $O(n)$  work which is often impractical.

None of the existing approaches support concurrent operations on arrays. O' Neill suggests having a lock for each element of the array. However, when many threads contend for the same element, this would serialize accesses and make accesses not  $O(1)$ . Additionally, per-element locks add significant memory overhead. Separation logic could be used to parallelize divide and conquer algorithms on arrays however they would severely restrict the way functional arrays can be used.

### 1.2 Our Approach

Our arrays support 3 methods: tabulate, get, and set. The dynamics are functional, in particular set returns a new array and leaves the original array intact.

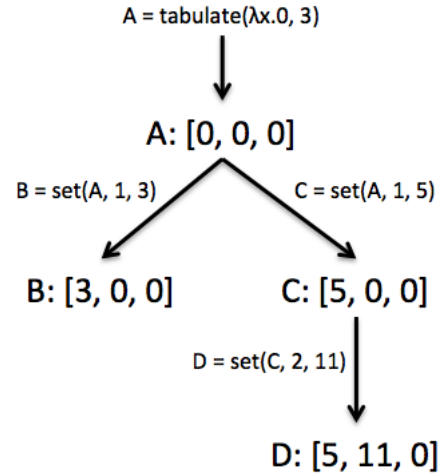


Figure 1. Example usage of functional arrays

In figure ??, A and C are *interior* arrays and B and D are *leaf* arrays. Calling get and set on leaf arrays costs  $O(1)$  but calling get and set on interior arrays is more expensive. In the cost dynamics, we thread a store that keeps track of whether an array is an interior or leaf array.

We also give and prove fork join cost dynamics for our arrays. The dynamics for arrays used concurrently are non trivial because one thread might call get on a leaf array while another thread calls set. To deal with this, we compute the work of each fork separately, and then charge additional work at the join point if an array is used in both forks.

In our implementation of a functional array, we keep an array of values (multiple functional arrays might point to the same array). For each index of the array, we also keep a history of values that

were previously at that index. To access an element in an interior array, we need to search through the history of values.

We give a wait-free concurrent implementation that uses careful synchronization. We prove linearizability of the implementation and use that to prove that arrays work correctly when used concurrently. Finally, we give a provable implementation for our cost dynamics.

## 2. Dynamics

We use a standard applicative-order language defined as follows:

$$e = x \mid c \mid \lambda x.e \mid e_1 e_2 \mid e_1 || e_2 \mid \text{if } e_1 \text{ then } e_2 \text{ else } e_3$$

$L$  contains the usual arithmetic types, such as the natural numbers, and numerical operations such as sums and products. Our dynamics are given by the judgement  $e \Downarrow v$  which means that expression  $e$  evaluates to  $v$ . Dynamics for function application and fork-join are given as examples.

$$\frac{e_1 \Downarrow v_1 \quad e_2 \Downarrow v_2}{e_1 || e_2 \Downarrow (v_1, v_2)} \text{ (fork-join)}$$

$$\frac{f \Downarrow \lambda x.e \quad y \Downarrow v}{f y \Downarrow [v/x]e} \text{ (function-app)}$$

The language also includes 3 functions for working with arrays: `tabulate`, `get`, and `set`. The dynamics for `tabulate` is given below:

$$\frac{f(1) \Downarrow v_1, \dots, f(n) \Downarrow v_n}{\text{tabulate } f \ n \Downarrow [1 \mapsto v_1, \dots, n \mapsto v_n]}$$

`get( $A, i$ )` returns the value at the  $i$ th index of  $A$ . `set( $A, i, v$ )` returns a new array where the value at index  $i$  is  $v$  and the value at all other indices is the same as in  $A$ . The evaluation of a simple program in our language is shown below:

```
(λA. get(A, 2) + get(A, 3))set(tabulate(λi. i, 5), 2, 10)
(λA. get(A, 2) + get(A, 3))set([1, 2, 3, 4, 5], 2, 10)
(λA. get(A, 2) + get(A, 3))[1, 10, 3, 4, 5]
get([1, 10, 3, 4, 5], 2) + get([1, 10, 3, 4, 5], 3)
10 + 3
13
```

## 3. Cost Dynamics

The dynamics given in the previous section are functional. However, the costs of `get` and `set` depend on whether the array is a leaf array or an interior array. To deal with this, we thread a store through the cost dynamics to keep track of whether each array is a leaf or interior array.

The cost dynamics of fork-join pose an additional challenge. An array can be used concurrently by multiple threads and the costs could depend on the order in which the instructions are interleaved. Figure ?? shows 3 ways a leaf array  $A$  might be used. In particular, one side of the fork might call `get` on  $A$  while the other side calls `set` (see the center panel). The calls to `get` on the left branch would be  $O(1)$  if and only if they execute before the call to `set` on the right branch.

We cannot make any assumptions about how the instructions in different branches of a fork are interleaved, so we give the worst case time complexity over all possible interleavings. We first compute the work of each fork separately. Then for arrays that are used in both forks, we add additional work at the join point. To compute the additional work at the join point, we keep track of the number of gets that cost  $O(1)$  on each side of the fork. If the array was modified on the other side of the fork, then the modification might have happened before the `get`, so we need to charge additional work for the `get`.

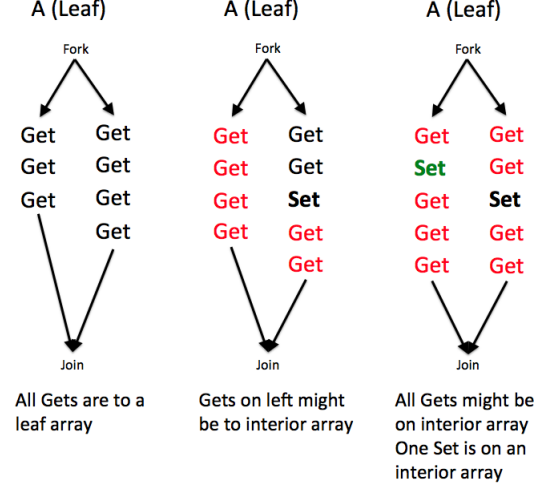


Figure 2. Different ways an array might be used in a fork join

Our cost dynamics are defined by the following judgement, where  $\delta$  is the store,  $e$  is the expression,  $\delta'$  is the new store,  $v$  is the value returned by evaluating  $e$ ,  $w$  is the work, and  $s$  is the span.

$$\delta; f \text{ args} \Downarrow \delta'; v; w; s$$

We give the cost dynamics for expressions in the language: (insert cost dynamics for standard functions later)

Arrays are represented by  $(l, V)$  where  $l$  is a label that uniquely identifies the array, and  $V$  contains the values in the array.  $V[i]$  is the  $i$ th value in the array. If the `set` method has been applied on an array, the array is a leaf array (represented by  $-$ ), otherwise the array is an interior array (represented by  $+$ ).

Array methods have different costs depending on whether the array is a leaf or interior array. Let  $C(l)$  be the number of `get` operations applied on array  $A = (l, V)$  that had  $O(1)$  work. We thread a store  $\delta$  through each operation.  $\delta$  maps array labels  $l$  to  $(+/-, c)$ , where  $c = C(l)$ . Let  $f(n)$  be the work of calling `get` on an interior array and  $g(n)$  be the work of calling `set` on an interior array.

The cost dynamics for the methods on a functional array of size  $n$  are:

$$\frac{l = \text{new label} \quad f(1) \Downarrow v_1; w_1; s_1; \dots; f(n) \Downarrow v_n; w_n; s_n}{\delta; \text{tabulate } f \ n \Downarrow \delta[l \mapsto (+, 0)]; (l, \bigcup_{i=1}^n [i \mapsto v_i]); 1 + \sum w_i; 1 + \max s_i}$$

$$\frac{A = (l, V) \quad \delta[l \mapsto (+, c)] \quad l' = \text{new label}}{\delta; \text{set } A \ i \ v \Downarrow \delta[l \mapsto (-, c), l' \mapsto (+, 0)]; (l', V[i \mapsto v]); 1; 1} \text{ (set-leaf)}$$

$$\frac{A = (l, V) \quad \delta[l \mapsto (-, c)] \quad l' = \text{new label}}{\delta; \text{set } A \ i \ v \Downarrow \delta[l \mapsto (-, c), l' \mapsto (+, 0)]; (l', V[i \mapsto v]); g(n); 1} \text{ (set-interior)}$$

$$\frac{A = (l, V) \quad \delta[l \mapsto (+, c)]}{\delta; \text{get } A \ i \Downarrow \delta[l \mapsto (+, c+1)]; V[i]; 1; 1} \text{ (get-leaf)}$$

$$\frac{A = (l, V) \quad \delta[l \mapsto (-, c)]}{\delta; \text{get } A \ i \Downarrow \delta; V[i]; f(n); 1} \text{ (get-interior)}$$

`set` creates a new label and array and modifies the store to indicate that the newly created array is a leaf array, and the array `set` was called on has become an interior array. The work for `set` is higher for interior arrays. `get` returns the value at the desired index, if the array is a leaf then the operation costs  $O(1)$  but we increment the counter of  $O(1)$  gets in the store.

The fork-join cost semantics are the most interesting. We give the cost dynamic using 2 helper functions. Let  $L(\delta)$  denote the set

of labels in the store  $\delta$ . Simplify the fork join cost dynamics

def NEW-MAP:  $L(\delta_1) \cup L(\delta_2) \rightarrow (+/-, \text{int}) =$

Case  $\delta[l \mapsto (-, c)] : (-, c)$

Case  $\delta[l \mapsto (+, c)] :$

Case  $\delta_1[l \mapsto (+, c + x)], \delta_2[l \mapsto (+, c + y)] : (+, c + x + y)$

Case  $\delta_1[l \mapsto (+, c + x)], \delta_2[l \mapsto (-, c + y)] : (-, c + y)$

Case  $\delta_1[l \mapsto (-, c + x)], \delta_2[l \mapsto (+, c + y)] : (-, c + x)$

Case  $\delta_1[l \mapsto (-, c + x)], \delta_2[l \mapsto (-, c + y)] : (-, c)$

Else:

Case  $\delta_1[l \mapsto (s, c)] : (s, c)$

Case  $\delta_2[l \mapsto (s, c)] : (s, c)$

def EXTRA-WORK:  $L(\delta_1) \cup L(\delta_2) \rightarrow \text{int} =$

Case  $\delta[l \mapsto (+, c)] :$

Case  $\delta_1[l \mapsto (+, c + x)], \delta_2[l \mapsto (+, c + y)] : 0$

Case  $\delta_1[l \mapsto (+, c + x)], \delta_2[l \mapsto (-, c + y)] : xf(n)$

Case  $\delta_1[l \mapsto (-, c + x)], \delta_2[l \mapsto (+, c + y)] : yf(n)$

Case  $\delta_1[l \mapsto (-, c + x)], \delta_2[l \mapsto (-, c + y)] : (x + y)f(n) + g$

Else: 0

$$\delta' = \bigcup_{l \in L(\delta_1) \cup L(\delta_2)} \text{NEW-MAP}(l)$$

$$w' = \sum_{l \in L(\delta_1) \cup L(\delta_2)} \text{EXTRA-WORK}(l)$$

Then, the fork-join cost semantics are:

$$\delta; e_1 || e_2 \Downarrow \delta'; (v_1, v_2); 1 + w_1 + w_2 + w'; 1 + \max(s_1, s_2)$$

## 4. Approach

An array  $A$  is NEW if the set function has not been called with  $A$  as an argument, otherwise  $A$  is OLD. While the dynamics of our arrays are functional, the costs are different for OLD and NEW arrays.




Figure 3. New functional array filled with 0s

Suppose that  $A$  is a NEW functional array (see figure ??).  $A$  has a version number  $V$ , and a pointer to an ArrayData object  $AD$ .  $AD$  keeps a regular (mutable) array of values, which corresponds to the values in  $A$ .  $AD$  has a version number, which is the same as  $A$ 's ( $V$ ) to indicate that  $A$  is NEW. For each element of the array in  $AD$ , keep a log of values that used to be at that index. The logs are initially empty.

Get on NEW arrays: to get the  $i^{\text{th}}$  element in  $A$ , simply access  $\text{array}[i]$  in  $AD$ .

Set on NEW arrays: Suppose that  $A[i]$  is  $v_{\text{old}}$  and we want to do  $B = \text{set}(A, i, v_{\text{new}})$  (see figure ??). Add an entry  $(V, v_{\text{old}})$  to the log at index  $i$ , expressing that  $A[i]$  was  $v_{\text{old}}$  at version  $V$ . Increment the version of  $AD$  to  $V + 1$ . Set  $\text{array}[i]$  in  $AD$  to  $v_{\text{new}}$ . Create a new functional array  $B$ , with version  $V + 1$ , which points to  $AD$ . Notice that both  $A$  and  $B$  point to  $AD$ . However, the version of  $A$  is  $V$  and the version of  $AD$  is  $V + 1$ , which indicates that  $A$  is OLD. The version of  $B$  is  $V + 1$  which indicates that  $B$  is NEW.




Figure 4.  $B = \text{set}(A, 3, 5)$  changes the array in  $AD$  and adds a log entry.

Get on OLD array: Suppose we do  $\text{get}(A, i)$  where  $A$  is OLD and has version  $V$ . We binary search the log at index  $i$  to find what value was stored in the array at version  $V$ .




Figure 5.  $D = \text{set}(A, 1, 9)$  copies  $AD$  because  $A$  is OLD

Suppose we do  $D = \text{set}(A, i, v)$  where  $A$  is old (see figure ??). Create a new ArrayData object  $AD'$  (with a new array). Copy the values from the array in  $AD$  to the array in  $AD'$ , and then set  $\text{array}[i]$  to  $v$  in  $AD'$ .

The logs are stored in a doubling array which doubles in size when full. This guarantees amortized  $O(1)$  insertion of a log entry, and allows us to binary search with  $O(\log m)$  work, where  $m$  is the number of log entries.

Suppose that the size of the array is  $n$ . Every  $n$  times an ArrayData object is modified, we create a new ArrayData object, and copy the values over to the new ArrayData object. This ensures that  $m \leq n$  i.e. we don't have too many log entries in an ArrayData object. Copying is amortized  $O(1)$ .

It follows that in the RAM model the work of set and get is  $O(1)$  for a NEW array. Get in an OLD array involves a binary search and has work  $O(\log n)$ . Set in an OLD array involves copying the array, and has work  $O(n)$ .

## 5. Concurrent Implementation

We implement the functional array functions using an imperative target language.

### 5.1 Concurrency Model

We assume a sequential consistency model of concurrency. That is, when analyzing the effects of a program, we assume that at each time step exactly one pending instruction is executed. Consider the execution of a program  $P$ . We label the time steps in the execution from 1 to  $t$ .

**Definition 5.1.** Consider an object  $A$ . Suppose that the program invokes methods  $M_1, M_2, \dots, M_n$  on  $A$ . Consider arbitrary  $i$ , and suppose that the instructions in  $M_i$  were executed at times  $s_1, s_2, \dots, s_m$ . We say that  $A$  is *linearizable* if the effect of the program is

as if  $M_i$  executed atomically at some time  $s_j$ , and this holds for all  $M_i$ .

**Definition 5.2.** Suppose that first instruction in function call  $f$  was executed at time  $s_1$  and the last instruction at time  $e_1$ . Suppose that first instruction in function call  $g$  was executed at time  $s_2$  and the last instruction at time  $e_2$ . We say the function calls do not *overlap* if the intervals  $[s_1, e_1]$  and  $[s_2, e_2]$  do not overlap.

Our programs have access to the link load store conditional (LLSC) function. LLSC takes an (address, old value, new value) tuple as argument. LLSC first checks that the value at address is the same as old value, if it is different then LLSC returns false. Otherwise LLSC atomically stores new value at address and returns true, however if the value at address was modified between the load and store, LLSC returns false.

## 5.2 Push Arrays

We use an auxiliary data structure, *PushArrays*, to store log entries. PushArrays are initially empty and support 3 methods. Suppose we have a PushArray  $A$ .  $A.push(e)$  inserts entry  $e$ ,  $A.size()$  returns the number of entries inserted,  $A.get(i)$  returns the  $i^{\text{th}}$  entry inserted. Note that  $get$  is only defined between indices 0 and  $A.size() - 1$ . All operations are amortized  $O(1)$ .

PushArrays can be used semi-concurrently. At most one thread can execute  $push$  at any time, but multiple threads can call  $size$  and  $get$ . In a PushArray  $A$ ,  $A.size$  stores the number of entries and  $A.data$  references an array that stores the actual entries. Pseudocode for  $A.push(e)$  is shown below.

```
if (isFull(A.data)) {
    newData = new Array(A.data.capacity * 2);
    copyValues(from = A.data, to = newData);
    A.data = newData;
}
A.data[A.size] = e;
A.size += 1;
```

**Theorem 5.1.** Given that different calls to  $push$  (on the same PushArray) do not overlap, PushArrays are linearizable.

*Proof.* The linearization point of  $push$  is  $A.size + 1$ , and the linearization point of the other methods is their single instruction. In  $push$ , expanding the capacity of  $A.data$  when it is full does not interfere with accessing  $A.data$  because we copy values to the new array before pointing  $A.data$  to the new array. The effects of adding  $e$  are only observable after we increase  $size$ , since, by the specification of PushArrays, the result of calling  $A.get(i)$  where  $i \geq A.size()$  is not defined.  $\square$

## 5.3 Functional Array Implementation

Suppose we have a FunctionalArray  $A$ .  $A$  points to an ArrayData object which is referenced by  $A.data$ . The (imperative) array in  $A.data$  is referenced by  $A.data.values$ , and the log entries corresponding to the  $i^{\text{th}}$  index are stored in a PushArray  $A.data.undo.lists[i]$ . We first explain the implementation of  $A.set(pos, val)$ .

We use a link load store conditional to ensure that only 1 thread can modify an ArrayData object at any time. If  $m$  threads try to modify  $A.data$ , then  $m - 1$  threads will fail the LLSC (returning false). Instead of modifying  $A.data$ , these threads will create a new ArrayData object, and will copy over values from  $A.data$ , as shown below.

```
FunctionalArray newA;
newA.version = A.version + 1;
if (!LLSC(&A.data.version, A.version,
```

```
newA.version) OR
A.version % A.size == 0) {
    // Create a new ArrayData object AD
    // Forall i = 1:n, AD.values[i] = get(A, i)
    AD.values[pos] = val;
    newA.data = AD;
    return newA;
}
```

At most one of the threads trying to modify the ArrayData object will get past the LLSC, will insert a log entry and then modify  $A.data.values$ . The ordering of inserting a log entry and modifying  $A.data.values$  matters in the proof of  $get$ . The pseudocode for this step is shown below.

```
newA.data = A.data;
old_value = newA.data.values[pos];
undo_list = newA.data.undo.lists[pos];
undo_list.push((A.version, old_value));
newA.data.values[pos] = val;
return newA;
```

Next we explain the implementation of  $A.get(pos)$ .

```
guess_val = A.data.values[pos];
if (A.version == A.data.version) {
    return guess_val;
}

undo_list = this.data.undo.lists[pos];
if (undo_list.size() == 0) {
    return guess_val;
}

upper = undo_list.size() - 1;
if (undo_list[upper].version < A.version) {
    return guess_val;
}
```

The last step is to binary search  $undo\_list$ , between the indices 0 and  $upper$  (inclusive of 0 and  $upper$ ). We find the log entry  $L$  with the smallest version  $X$  such that  $A.version \leq X$ . We return  $L.value$ .

The implementation of  $tabulate$  is straightforward - just allocate a new array and fill the array with the required values.

## 5.4 Proof of Correctness

**Theorem 5.2.** Functional Arrays are linearizable.

*Proof.* The linearization point of  $get$  is when we compare if  $A.version$  and  $AD.version$  are the same. The linearization point of  $set$  is the LLSC. Any instruction in  $tabulate$  can be its linearization point - we omit discussion of  $tabulate$  because the details are uninteresting.

First, we prove the linearization point of  $get$ . (fill this out)

Next, we prove the linearization point of  $set$ . (fill this out)  $\square$

Linearizability allows us to assume that array operations happen atomically. Note that all non array operations in our source language are executed in a single time step, so any well formed program is linearizable. In fact, since linearizability composes, we can show that more complicated languages that use our arrays are also linearizable.

In a previous section, we gave the evaluational (big step) dynamics of functional arrays. Alternatively, we could give the non-deterministic structural (little step) dynamics which describes the steps in the execution of a program in the source language. At each

fork-join, if both sides of the fork can take a step, we could take a step on either side.

**Theorem 5.3.** *The non-deterministic structural dynamics and the evaluational dynamics produce the same result in the source language.*

*Proof.* This is a standard implicit parallelism type theorem (ref: PFPL) so we omit the details.  $\square$

Now, to show that our implementation correctly implements the source language, it suffices to show that any valid execution of our implementation produces the same result as a little step transition sequence in the source language (since all transition sequences produce the same value as the evaluational dynamics).

**Theorem 5.4.** *For any program  $P$  in the source language, executing our implementation  $I$  of  $P$  produces the same result as some transition sequence of  $P$ . Note that  $I$  is a program in the target language.*

*Proof.* We prove the claim by induction. We focus on the part of the proof dealing with arrays and give a sketch of the rest.

Our inductive hypothesis is that after  $t$  time steps in an arbitrary execution of  $I$ , there exists a transition sequence  $S$  s.t.:

1.  $I$  becomes  $I'$  after this particular execution,  $S$  takes  $P$  to  $P'$
2.  $I'$  and  $P'$  are equivalent in the following sense. Any non array value is identical in  $I'$  and  $P'$ . For any abstract array  $A$  in  $P'$  there exists a corresponding representation  $A_I$  s.t. the following holds. Consider arbitrary index  $i$  and suppose that  $A[i]$  has value  $val$ . Suppose that  $A_I$  has version  $V$  and is pointing to an ArrayData object  $AD$ . If there exists some version in  $AD.logs[i]$  which is greater than  $V$  then there exists a log entry  $(v', val)$  where  $v'$  is the smallest version in  $AD.logs[i]$  that is  $\geq V$ . If no version in  $AD.logs[i]$  is greater than  $V$ , then  $AD.values[i]$  is  $val$ .

The base case holds since  $P$  and  $I$  are initially identical. For the inductive step, when  $I'$  takes a step, we take the corresponding step in  $P$  to get a new transition sequence  $S'$ . We know that such a step exists because  $I'$  and  $P'$  are equivalent. If the step is not an array function then the IH holds since the step will be executed in the same way in the source and target languages.

Suppose that the next step is tabulate. By the IH, tabulate is called with the same arguments  $n$  and  $f$  in the source language and the implementation. In the source language, tabulate returns an array  $A$  with  $A[i] = f(i)$  for all  $1 \leq i \leq n$ . In the implementation, tabulate returns an array that points to a new ArrayData object  $AD$  with empty logs and  $AD.values[i] = f(i)$  for all  $1 \leq i \leq n$ . So the inductive step holds.

Suppose that the next step is get. By the IH, get is called on equivalent arrays and the same index  $i$  in the source language and implementation. It is easy to trace through the implementation of get and see that get returns the same value in the source language and implementation.

Suppose that the next step is set. Suppose that in the target language we execute  $set(A_I, i, v)$ . Then by IH, in the source language the corresponding step is  $set(A, i, v)$  where  $A$  and  $A_I$  are equivalent. Denote  $set(A, i, v)$  by  $B$  and  $set(A_I, i, v)$  by  $B_I$ . We need to show that  $B$  and  $B_I$  are equivalent, and all arrays that were previously equivalent are still equivalent.

Case 1: we create a new ArrayData object  $AD$  in set. In this case, for all  $j \neq i$  we set  $AD.values[j]$  to  $get(A_I, j)$  which by the IH is the same as  $A[j]$ . Since we only modified index  $i$ ,  $j \neq i \Rightarrow A[j] = B[j]$ . Also, we set  $AD.values[i]$  to  $v$ , so  $AD.values[i] = B[i]$ . Since all logs are empty, this means that  $B$  and  $B_I$  are equivalent. Since we create a new ArrayData object,

we do not modify any other arrays, so arrays that were previously equivalent are still equivalent.

Case 2: we modify the ArrayData object  $AD$  that  $A_I$  points to. This means that  $A.version = AD.version$ . We can show that all log entries in  $AD$  have version  $< AD.version$ , and are in increasing order. Then the IH implies that before modifying  $AD$ ,  $AD.values[j] = A[j]$  for all indices  $j$ . So after setting  $AD.values[i]$  to  $v$ ,  $AD.values[j] = B[j]$  for all  $j$ , which means  $B_I$  and  $B$  are equivalent.

Consider arbitrary abstract array  $C$  in the source language and corresponding representation  $C_I$  in the implementation with  $C_I \neq B_I$ . If  $C_I$  did not point to  $AD$  then  $C_I$  remains equivalent to  $C$ . Regardless, we note that for all indices  $j \neq i$ ,  $C_I$  remains the same and is equivalent to  $C$ . At index  $i$ , if the value at  $C[i]$  was previously stored in a log entry, then  $C_I$  and  $C$  remain equivalent since log entries are added in increasing version numbers. If the value at  $C[i]$  was stored in  $AD[i]$  then we add a log entry at index  $i$  so  $C_I$  and  $C$  remain equivalent.  $\square$

## 6. Machine Model

Our target language is a Java-like language augmented with tabulate, fork-join, lambda functions, and link load store conditional (LLSC). Fork-join spawns two threads, one thread executing each sub-expression. Link load store conditional (LLSC) takes an (address, old value, new value) tuple as argument. LLSC first checks that the value at address is the same as old value, if it is different then LLSC returns false. Otherwise LLSC atomically stores new value at address and returns true, however if the value at address was modified between the load and store, LLSC returns false.

Our target language is run on a P processor machine. In a single time step the machine takes at most P runnable instructions and executes them. The effect of executing the ( $\leq P$ ) instructions in a time step is the same as some (arbitrary) sequential ordering of the instructions. We assume that the machine does not reorder instructions in the target language. In reality, languages like Java have very relaxed consistency models, and we deal with this in our real implementation using memory fences.

The link load store conditional architecture deserves special mention. Call an LLSC operation pending if the load operation has completed but the store operation has not. Multiple threads can load the value at a particular address and compare it with old value in a single time step. Only a single store operation to a particular address can be executed at a time step. However, after the store operation, in each time step the bus arbiter notifies all pending LLSC operations to the same memory address that they have failed. This assumption is reasonable because all LLSC instructions are typically processed by the bus arbiter, so the bus arbiter can keep track of pending LLSC instructions to each memory address.